

**Question 1. (25 points)** In the Laffer data (*Data/Laffer.csv*, first introduced in the accompanying prediction tutorial), there is one country with a high tax revenue that is an outlier.

In the tutorial, PSIS and WAIC were used to measure the importance of this outlier in the models that were fit.

First, re-create and run the models provided as part of the prediction tutorial.

```
In [43]: df = pd.read_csv("Data/Laffer.csv", delimiter=",")

df = pd.read_csv("Data/Laffer.csv", header=0)
# standardizing tax rate and tax revenue
tax_rate_std = standardize(df.tax_rate)
tax_revenue_std = standardize(df.tax_revenue)

# first-order linear model
with pm.Model() as m_ex4_lin:

    # prior on intercept for expected tax revenue: mean = 0, std = 0.2
    a = pm.Normal("a", 0, 0.2)

    # prior on coefficient for tax rate: mean=0, std=0.5
    b = pm.Normal("b", 0, 0.5)

    # expected tax revenue (using linear model)
    mu = pm.Deterministic("mu", a + b * tax_rate_std)

    # standard deviation for tax revenue
    sigma = pm.Exponential("sigma", 1)

    # data distribution for tax revenue
    R = pm.Normal("tax_revenue", mu, sigma, observed=tax_revenue_std)

    # using Hamiltonian Monte Carlo MCMC
    idata_ex4_lin = pm.sample(idata_kwargs={"log_likelihood": True})

# quadratic model
with pm.Model() as m_ex4_quad:
    # prior on intercept for expected tax revenue: mean = 0, std = 0.2
    a = pm.Normal("a", 0, 0.2)

    # prior on coefficient for tax rate: mean=0, std=0.5
    b = pm.Normal("b", 0, 0.5)

    # prior on coefficient for tax rate^2: mean=0, std=0.5
    b2 = pm.Normal("b2", 0, 0.5)

    # expected tax revenue (using polynomial model)
    mu = pm.Deterministic("mu", a + b * tax_rate_std + b2 * tax_rate_std**2)

    # standard deviation for tax revenue
    sigma = pm.Exponential("sigma", 1)
```

```

# data distribution for tax revenue
R = pm.Normal("tax_revenue", mu, sigma, observed=tax_revenue_std)

# using Hamiltonian Monte Carlo MCMC
idata_ex4_quad = pm.sample(idata_kwargs={"log_likelihood": True})

```

```

Auto-assigning NUTS sampler...
Initializing NUTS using jitter+adapt_diag...
Sequential sampling (2 chains in 1 job)
NUTS: [a, b, sigma]

```

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

```

Sampling 2 chains for 1_000 tune and 1_000 draw iterations (2_000 + 2_000 draws total) took 5 seconds.
We recommend running at least 4 chains for robust computation of convergence diagnostics
Auto-assigning NUTS sampler...
Initializing NUTS using jitter+adapt_diag...
Sequential sampling (2 chains in 1 job)
NUTS: [a, b, b2, sigma]

```

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

```

Sampling 2 chains for 1_000 tune and 1_000 draw iterations (2_000 + 2_000 draws total) took 7 seconds.
We recommend running at least 4 chains for robust computation of convergence diagnostics

```

Create a new **cubic** (3rd-degree polynomial) version of the model.

```
In [44]: # cubic model
with pm.Model() as m_ex4_cubic:
    # prior on intercept for expected tax revenue: mean = 0, std = 0.2
    a = pm.Normal("a", 0, 0.2)

    # prior on coefficient for tax rate: mean=0, std=0.5
    b = pm.Normal("b", 0, 0.5)

    # prior on coefficient for tax rate^2: mean=0, std=0.5
    b2 = pm.Normal("b2", 0, 0.5)

    # prior on coefficient for tax rate^3: mean=0, std=0.5
    b3 = pm.Normal("b3", 0, 0.5)

    # expected tax revenue (using polynomial model)
    mu = pm.Deterministic("mu", a + b * tax_rate_std + b2 * tax_rate_std**2 + b3 * tax_rate_std**3)

    # standard deviation for tax revenue
    sigma = pm.Exponential("sigma", 1)

    # data distribution for tax revenue
    R = pm.Normal("tax_revenue", mu, sigma, observed=tax_revenue_std)

    # using Hamiltonian Monte Carlo MCMC
    idata_ex4_cubic = pm.sample(idata_kwargs={"log_likelihood": True})
```

```
Auto-assigning NUTS sampler...
Initializing NUTS using jitter+adapt_diag..
Sequential sampling (2 chains in 1 job)
NUTS: [a, b, b2, b3, sigma]
```

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

Sampling 2 chains for 1\_000 tune and 1\_000 draw iterations (2\_000 + 2\_000 draws total) took 15 seconds.  
We recommend running at least 4 chains for robust computation of convergence diagnostics

Create a new (first-order) **linear** version of the model using robust regression with a Student's t distribution for the data distribution. Use  $\nu = 2$  for your Student's t distribution.

```
In [45]: # first-order linear model using robust regression with a Student's t distribution for the data
with pm.Model() as m_ex4_lin_robust:
    # Prior on intercept for expected tax revenue: mean = 0, std = 0.2
    a = pm.Normal("a", 0, 0.2)

    # Prior on coefficient for tax rate: mean=0, std=0.5
    b = pm.Normal("b", 0, 0.5)

    # Expected tax revenue (using linear model)
    mu = pm.Deterministic("mu", a + b * tax_rate_std)

    # Standard deviation for tax revenue
    sigma = pm.Exponential("sigma", 1)

    # Degrees of freedom
    nu = 2
    R = pm.StudentT("tax_revenue", nu=nu, mu=mu, sigma=sigma, observed=tax_revenue_std)

    # using Hamiltonian Monte Carlo MCMC
    idata_ex4_lin_robust = pm.sample(idata_kwargs={"log_likelihood": True})

# Set up the PSIS comparison
model_dict = {
    "m_ex4_lin": idata_ex4_lin,
    "m_ex4_quad": idata_ex4_quad,
    "m_ex4_cubic": idata_ex4_cubic,
    "m_ex4_lin_robust": idata_ex4_lin_robust
}

# Execute the PSIS comparison
compare_df_psis = az.compare(
    compare_dict=model_dict,
    ic="loo",
    scale="deviance"
)
```

```
Auto-assigning NUTS sampler...
Initializing NUTS using jitter+adapt_diag..
Sequential sampling (2 chains in 1 job)
NUTS: [a, b, sigma]
```

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

```
Sampling 2 chains for 1_000 tune and 1_000 draw iterations (2_000 + 2_000 draws total) took 5 seconds.
We recommend running at least 4 chains for robust computation of convergence diagnostics
/opt/conda/envs/fnds/lib/python3.9/site-packages/arviz/stats/stats.py:789: UserWarning: Estimated shape
  warnings.warn(
/opt/conda/envs/fnds/lib/python3.9/site-packages/arviz/stats/stats.py:789: UserWarning: Estimated shape
  warnings.warn(
/opt/conda/envs/fnds/lib/python3.9/site-packages/arviz/stats/stats.py:789: UserWarning: Estimated shape
  warnings.warn(
```

Of the 4 models, which one has the best predictive performance according to PSIS? **Limit your response to no more than 2 sentences.**

The linear version using robust regression has the highest rank (0) out of the 4 models. The value for elpd\_loo is ~73.84 which indicates a better predictive performance compared to ~86 for the others, it generalizes data better.



Based on these results, what impact does the outlier point appear to be having on the curved relationship between tax rate and tax revenue? **Limit your response to no more than 5 sentences.**

**Note:** *You will need to quantify the influence of the outlier data point in your predictive models and comment on this influence in order to receive full credit.*

```
In [46]: print(compare_df_psis)
```

	rank	elpd_loo	p_loo	elpd_diff	weight \
m_ex4_lin_robust	0	73.664599	3.123187	0.000000	7.789287e-01
m_ex4_quad	1	85.623707	5.161043	11.959108	2.210713e-01
m_ex4_lin	2	87.286892	5.070848	13.622293	4.906627e-16
m_ex4_cubic	3	87.353628	5.968676	13.689029	0.000000e+00

  

	se	dse	warning	scale
m_ex4_lin_robust	12.261320	0.000000	False	deviance
m_ex4_quad	19.265196	11.671712	True	deviance
m_ex4_lin	18.992424	11.326380	True	deviance
m_ex4_cubic	19.592405	11.809895	True	deviance

For the quadratic linear (non-robust), and cubic models, the elpd\_loo value is ~86 while the linear (robust) has an elpd\_loo value of ~73.84. The linear robust model is influenced less by the outlier value in our data, making it a better fitting model. The others elpd\_loo values indicate overfitting at a value of ~86, in addition to their se values being ~18 or 19, which is higher than the se value for linear (robust) se at a value of ~12.28.





**Question 2. (15 points)** Reconsider the urban fox analysis (*Data/foxes.csv*) from the previous homework assignment (HW 3).

On the basis of PSIS and WAIC scores, which combination of variables best predicts body weight (*W*, weight)? **Limit your response to no more than 3 sentences.**

*Note: For this exercise, you will end up creating 7 models.*

```
In [47]: df = pd.read_csv(
        "Data/foxes.csv",
        sep=',',
        header=0
    )
    # prior predictive simulation
    a_std = standardize(df.area)
    g_std = standardize(df.group)
    f_std = standardize(df.avgfood)
    gs_std = standardize(df.groupsize)
    w_std = standardize(df.weight)
    # Store models in a structure
    models = {}
    # Area model
    with pm.Model() as m_AREA:
        a = pm.Normal("a", 0.1, 0.2)
        b_a = pm.Normal("b_a", 0.1, 0.5)
        mu = pm.Deterministic("mu", a + b_a * a_std)
        sigma = pm.Exponential("sigma", 1)
        div = pm.Normal("div", mu, sigma, observed=w_std)
        # Using the quap method for inference (classical approach)
        idata_AREA = pm.sample(idata_kwargs={"log_likelihood": True})
    models["area"] = idata_AREA

    # Average food model
    with pm.Model() as m_GROUP:
        a = pm.Normal("a", 0.1, 0.2)
        b_f = pm.Normal("b_f", 0.1, 0.5)
        mu = pm.Deterministic("mu", a + b_f * f_std)
        sigma = pm.Exponential("sigma", 1)
        div = pm.Normal("div", mu, sigma, observed=w_std)
        idata_AVGFOOD = pm.sample(idata_kwargs={"log_likelihood": True})
    models["AVGFOOD"] = idata_AVGFOOD

    # Group size model
    with pm.Model() as m_GROUPSIZE:
        a = pm.Normal("a", 0.1, 0.2)
        b_gs = pm.Normal("b_gs", 0.1, 0.5)
        mu = pm.Deterministic("mu", a + b_gs * gs_std)
        sigma = pm.Exponential("sigma", 1)
        div = pm.Normal("div", mu, sigma, observed=w_std)
        idata_GROUPSIZE = pm.sample(idata_kwargs={"log_likelihood": True})
    models["GROUPSIZE"] = idata_GROUPSIZE
```

```

# Groupsize and area model
with pm.Model() as m_GROUPSIZE_AREA:
    a = pm.Normal("a", 0.1, 0.2)
    b_gs = pm.Normal("b_gs", 0.1, 0.5)
    b_a = pm.Normal("b_a", 0.1, 0.5)
    mu = pm.Deterministic("mu", a + b_gs * gs_std + b_a * a_std)
    sigma = pm.Exponential("sigma", 1)
    div = pm.Normal("div", mu, sigma, observed=w_std)
    idata_GROUPSIZE_AREA = pm.sample(idata_kwargs={"log_likelihood": True})
models["GROUPSIZE_AREA"] = idata_GROUPSIZE_AREA

# Average food and groupsize and area model
with pm.Model() as m_AVGFOOD_GROUPSIZE_AREA:
    a = pm.Normal("a", 0.1, 0.2)
    b_a = pm.Normal("b_a", 0.1, 0.5)
    b_f = pm.Normal("b_f", 0.1, 0.5)
    b_gs = pm.Normal("b_gs", 0.1, 0.5)
    mu = pm.Deterministic("mu", a + b_f * f_std + b_gs * gs_std + b_a * a_std)
    sigma = pm.Exponential("sigma", 1)
    div = pm.Normal("div", mu, sigma, observed=w_std)
    idata_AVGFOOD_GROUPSIZE_AREA = pm.sample(idata_kwargs={"log_likelihood": True})
models["AVGFOOD_GROUPSIZE_AREA"] = idata_AVGFOOD_GROUPSIZE_AREA

# Average food and groupsize model
with pm.Model() as m_AVGFOOD_GROUPSIZE:
    a = pm.Normal("a", 0.1, 0.2)
    b_f = pm.Normal("b_f", 0.1, 0.5)
    b_gs = pm.Normal("b_gs", 0.1, 0.5)
    mu = pm.Deterministic("mu", a + b_f * f_std + b_gs * gs_std)
    sigma = pm.Exponential("sigma", 1)
    div = pm.Normal("div", mu, sigma, observed=w_std)
    idata_AVGFOOD_GROUPSIZE = pm.sample(idata_kwargs={"log_likelihood": True})
models["AVGFOOD_GROUPSIZE"] = idata_AVGFOOD_GROUPSIZE

# Average food and area model
with pm.Model() as m_AVGFOOD_AREA:
    a = pm.Normal("a", 0.1, 0.2)
    b_f = pm.Normal("b_f", 0.1, 0.5)
    b_a = pm.Normal("b_a", 0.1, 0.5)
    mu = pm.Deterministic("mu", a + b_f * f_std + b_a * a_std)
    sigma = pm.Exponential("sigma", 1)
    div = pm.Normal("div", mu, sigma, observed=w_std)
    idata_AVGFOOD_AREA = pm.sample(idata_kwargs={"log_likelihood": True})
models["AVGFOOD_AREA"] = idata_AVGFOOD_AREA

# Execute the comparison between 7 models
# PSIS comparison
compare_df_psis_q2 = az.compare(
    compare_dict=models,
    ic="loo",
    scale="deviance"
)
# WAIC comparison

```

```

compare_df_haic_q2 = az.compare(
    compare_dict=models,
    ic="waic",
    scale="deviance"
)
print("PSIS:", compare_df_psis_q2)
print("WAIC:", compare_df_haic_q2)

```

Auto-assigning NUTS sampler...  
 Initializing NUTS using jitter+adapt\_diag...  
 Sequential sampling (2 chains in 1 job)  
 NUTS: [a, b\_a, sigma]

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

Sampling 2 chains for 1\_000 tune and 1\_000 draw iterations (2\_000 + 2\_000 draws total) took 4 seconds.  
 We recommend running at least 4 chains for robust computation of convergence diagnostics  
 Auto-assigning NUTS sampler...  
 Initializing NUTS using jitter+adapt\_diag...  
 Sequential sampling (2 chains in 1 job)  
 NUTS: [a, b\_f, sigma]

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

Sampling 2 chains for 1\_000 tune and 1\_000 draw iterations (2\_000 + 2\_000 draws total) took 4 seconds.  
 We recommend running at least 4 chains for robust computation of convergence diagnostics

```
Auto-assigning NUTS sampler...
Initializing NUTS using jitter+adapt_diag..
Sequential sampling (2 chains in 1 job)
NUTS: [a, b_gs, sigma]
```

```
<IPython.core.display.HTML object>
```

```
<IPython.core.display.HTML object>
```

```
<IPython.core.display.HTML object>
```

```
<IPython.core.display.HTML object>
```

```
Sampling 2 chains for 1_000 tune and 1_000 draw iterations (2_000 + 2_000 draws total) took 4 seconds.
We recommend running at least 4 chains for robust computation of convergence diagnostics
Auto-assigning NUTS sampler...
Initializing NUTS using jitter+adapt_diag..
Sequential sampling (2 chains in 1 job)
NUTS: [a, b_gs, b_a, sigma]
```

```
<IPython.core.display.HTML object>
```

```
<IPython.core.display.HTML object>
```

```
<IPython.core.display.HTML object>
```

```
<IPython.core.display.HTML object>
```

```
Sampling 2 chains for 1_000 tune and 1_000 draw iterations (2_000 + 2_000 draws total) took 7 seconds.
We recommend running at least 4 chains for robust computation of convergence diagnostics
Auto-assigning NUTS sampler...
Initializing NUTS using jitter+adapt_diag..
Sequential sampling (2 chains in 1 job)
NUTS: [a, b_a, b_f, b_gs, sigma]
```

```
<IPython.core.display.HTML object>
```

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

Sampling 2 chains for 1\_000 tune and 1\_000 draw iterations (2\_000 + 2\_000 draws total) took 11 seconds.  
We recommend running at least 4 chains for robust computation of convergence diagnostics  
Auto-assigning NUTS sampler..  
Initializing NUTS using jitter+adapt\_diag..  
Sequential sampling (2 chains in 1 job)  
NUTS: [a, b\_f, b\_gs, sigma]

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

Sampling 2 chains for 1\_000 tune and 1\_000 draw iterations (2\_000 + 2\_000 draws total) took 8 seconds.  
We recommend running at least 4 chains for robust computation of convergence diagnostics  
Auto-assigning NUTS sampler..  
Initializing NUTS using jitter+adapt\_diag..  
Sequential sampling (2 chains in 1 job)  
NUTS: [a, b\_f, b\_a, sigma]

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

Sampling 2 chains for 1\_000 tune and 1\_000 draw iterations (2\_000 + 2\_000 draws total) took 7 seconds.  
We recommend running at least 4 chains for robust computation of convergence diagnostics

PSIS:		rank	elpd_loo	p_loo	elpd_diff	weight \
AVGFOOD_GROUPSIZE_AREA	0	322.847144	4.439754	0.000000	2.885929e-15	
GROUPSIZE_AREA	1	323.534026	3.443261	0.686882	5.030876e-01	
AVGFOOD_GROUPSIZE	2	323.692270	3.459335	0.845126	4.969124e-01	
GROUPSIZE	3	330.740111	2.544852	7.892966	0.000000e+00	
AVGFOOD	4	333.415140	2.324560	10.567995	0.000000e+00	
area	5	333.515934	2.481409	10.668790	0.000000e+00	
AVGFOOD_AREA	6	334.676162	3.442385	11.829018	2.109424e-15	

	se	dse	warning	scale
AVGFOOD_GROUPSIZE_AREA	15.320499	0.000000	False	deviance
GROUPSIZE_AREA	14.995065	2.761174	False	deviance
AVGFOOD_GROUPSIZE	15.222967	3.308261	False	deviance
GROUPSIZE	14.018231	5.563166	False	deviance
AVGFOOD	13.265442	6.651596	False	deviance
area	13.221790	6.703250	False	deviance
AVGFOOD_AREA	13.333225	6.494722	False	deviance

WAIC:		rank	elpd_waic	p_waic	elpd_diff	weight \
AVGFOOD_GROUPSIZE_AREA	0	322.799783	4.416073	0.000000	2.704918e-15	
GROUPSIZE_AREA	1	323.508161	3.430329	0.708378	5.021318e-01	
AVGFOOD_GROUPSIZE	2	323.651538	3.438969	0.851755	4.978682e-01	
GROUPSIZE	3	330.718307	2.533950	7.918524	2.075288e-15	
AVGFOOD	4	333.400106	2.317043	10.600323	0.000000e+00	
area	5	333.497858	2.472371	10.698075	0.000000e+00	
AVGFOOD_AREA	6	334.639521	3.424064	11.839737	0.000000e+00	

	se	dse	warning	scale
AVGFOOD_GROUPSIZE_AREA	15.311801	0.000000	True	deviance
GROUPSIZE_AREA	14.991379	2.760849	True	deviance
AVGFOOD_GROUPSIZE	15.214082	3.303694	False	deviance
GROUPSIZE	14.013987	5.558503	False	deviance
AVGFOOD	13.263224	6.646099	False	deviance
area	13.219105	6.697645	False	deviance
AVGFOOD_AREA	13.327914	6.490053	False	deviance

```
/opt/conda/envs/fnds/lib/python3.9/site-packages/arviz/stats/stats.py:1632: UserWarning: For one or more
See http://arxiv.org/abs/1507.04544 for details
```

```
warnings.warn(
/opt/conda/envs/fnds/lib/python3.9/site-packages/arviz/stats/stats.py:1632: UserWarning: For one or more
See http://arxiv.org/abs/1507.04544 for details
warnings.warn(
```

The combination of variables Averagefood, Groupsize, and Area ranks highest (0) out of the 7 models. For elpd\_loo scores, this combination has the lowest score of ~322 compared to values greater than it for other

models, and `elpd_waic` has the highest value of  $\sim 4.27$ . The `elpd_diff` value and weight for this combination of variables also yields 0.00 for both, indicating a perfect fit.





Answer the following question based on your analysis above:

How would you interpret the estimates for the coefficients from the best scoring model in terms of their ability to predict weight? **Limit your response to no more than 10 sentences.**

*Hint: Answering this question requires an evaluation of the coefficient estimates which can be obtained by passing the inference data object for the best performing model to `arviz.summary()`.*

In [48]: # Write code here

```
az.summary(idata_AVGFOOD_GROUPSIZE_AREA, kind="stats")
```

```
Out[48]:
```

	mean	sd	hdi_5.5%	hdi_94.5%
a	0.019	0.079	-0.106	0.148
b_a	0.276	0.176	-0.006	0.551
b_f	0.295	0.217	-0.041	0.630
b_gs	-0.632	0.186	-0.926	-0.345
sigma	0.958	0.066	0.852	1.057
...	...	...	...	...
mu[111]	-0.086	0.090	-0.239	0.043
mu[112]	-0.086	0.090	-0.239	0.043
mu[113]	-0.311	0.202	-0.642	-0.003
mu[114]	-0.311	0.202	-0.642	-0.003
mu[115]	-0.311	0.202	-0.642	-0.003

[121 rows x 4 columns]

“a” represents the intercept value of the model with a mean of 0.016 ranging from -0.118 to 0.122 in a 89% HPDI, and has little influence on predicting weight. “b\_a” is the coefficient for area ranging from -0.026 to 0.522 so it could have little influence or it could have a moderate influence. “b\_f” is the coefficient of average food ranging from -0.061 to 0.639 so it could have little influence to a decent amount of impact on weight. “b\_gs” is the coefficient for group size and it ranges from -0.923 to -0.306 suggesting that weight is strongly impacted and decreasing as groupsize increases. Sigma suggests that in a range of 0.867 to 1.063 the model has room for uncertainty but fits the data.



**Question 3. (10 points)** Using the data in *Data/cherry\_blossoms.csv*:

In this problem, you will build predictive models of the relationship between the timing of cherry blossoms (*doy*) and March temperature in the same year (*temp*).

*Note:* \* Only include observations that have recorded values for *doy* and *temp* (i.e., exclude NaN values). \* The `pandas.DataFrame.dropna` method can be used to exclude NaN values.

Construct at least two different models to predict *doy* with *temp*.

```
In [ ]: df = pd.read_csv(
        "Data/cherry_blossoms.csv"
    )
    df_dropped = df.dropna(subset=["doy", "temp"])
    t_std = standardize(df_dropped["temp"])
    d_std = standardize(df_dropped["doy"])
    # first-order linear model
    with pm.Model() as cb_linear:
        a = pm.Normal("a", 0, 0.2)
        b = pm.Normal("b", 0, 0.5)
        mu = pm.Deterministic("mu", a + b * t_std)
        sigma = pm.Exponential("sigma", 1)
        R = pm.Normal("doy", mu, sigma, observed=d_std)
        idata_cb_linear = pm.sample(idata_kwargs={"log_likelihood": True})
    # quadratic model
    with pm.Model() as cb_quadratic:
        a = pm.Normal("a", 0, 0.2)
        b = pm.Normal("b", 0, 0.5)
        b_t = pm.Normal("b_t", 0, 0.5)
        mu = pm.Deterministic("mu", a + b * t_std + b_t * t_std**2)
        sigma = pm.Exponential("sigma", 1)
        R = pm.Normal("doy", mu, sigma, observed=d_std)
        idata_cb_quadratic = pm.sample(idata_kwargs={"log_likelihood": True})
```

```
Auto-assigning NUTS sampler...
Initializing NUTS using jitter+adapt_diag...
Sequential sampling (2 chains in 1 job)
NUTS: [a, b, sigma]
```

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

Sampling 2 chains for 1\_000 tune and 1\_000 draw iterations (2\_000 + 2\_000 draws total) took 4 seconds.  
We recommend running at least 4 chains for robust computation of convergence diagnostics  
Auto-assigning NUTS sampler..  
Initializing NUTS using jitter+adapt\_diag..  
Sequential sampling (2 chains in 1 job)  
NUTS: [a, b, b\_t, sigma]

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

<IPython.core.display.HTML object>

```
In [ ]: # Comparison using PSIS
        cherry_models = {
            "linear" : idata_cb_linear,
            "quadratic" : idata_cb_quadratic
        }
        compare_cherry_models = az.compare(cherry_models, ic="loo", scale="deviance")
        print(compare_cherry_models)
```

State which model is better at making predictions and explain what evidence leads you to this conclusion.  
**Limit your response to no more than 5 sentences.**

Linear is better at making predictions because it ranked higher (0) than quadratic (1). It has a lower elpd\_loo value than quadratic as well, as  $2148 < 2150$ . It is a better predictor than quadratic because the elpd\_diff value is 0.00000 whereas for quadratic, it is nonzero at a value of  $\sim 2.13$ .

```
In [ ]: grader.check("q3.1")
```

